

Fox ML Infrastructure LLC – Service Level Agreement (SLA)

Version 1.0 – Effective January 2026

Document Hierarchy: This document is provided for guidance only. In case of any conflict, COMMERCIAL_LICENSE.md is the authoritative and controlling document for all commercial licensing terms and obligations, and LICENSE (AGPL-3.0) controls for open-source licensing.

Commercial Licensees Only: This document applies only to customers with a valid commercial license. AGPL-3.0 users are not eligible for the policies, services, or terms described here.

SLA Opt-In Only: SLA applies only if expressly included in a signed Order Form or SOW. This is an internal policy summary provided for informational purposes only and does not create any SLA, guarantee, or obligation unless explicitly included in a signed agreement.

This Service Level Agreement (“SLA”) applies to customers purchasing the Enterprise Support tier from Fox ML Infrastructure LLC.

This SLA does not apply to Community or Standard Commercial license tiers.

1. Current State of Enterprise Support

Fox ML Infrastructure is in the process of finalizing:

- A dedicated enterprise support email inbox
- Scheduled support call availability
- Automated support ticket routing

These systems are currently being set up and will be fully active in early 2026.

Until then, Enterprise customers will receive direct communication instructions immediately upon purchase.

1.1 Pricing

Enterprise Support pricing: \$25,000-\$250,000/month based on SLA level and requirements. Pricing aligned with enterprise ML infrastructure support standards (Snowflake, Palantir, Databricks).

Tiered Pricing: - Business Support: \$25,000-\$50,000/month - Enterprise Support: \$60,000-\$120,000/month
- Premium Support: \$150,000-\$250,000/month

See [LEGAL/SUBSCRIPTIONS.md](#) for complete pricing information, tier definitions (response times, coverage, channels), and optional add-ons.

2. Definitions

Business Hours

8:00 AM - 6:00 PM US Central Time, Monday-Friday, excluding US holidays.

Incident

A reproducible failure of licensed Fox ML Infrastructure components to perform as documented.

Critical Incident

A failure preventing core system functionality with no available workaround.

Major Incident

A failure that impacts functionality but has a reasonable workaround.

Minor Incident

A non-critical bug or documentation clarification.

3. Response Time Guarantees

Fox ML Infrastructure provides the following response times for Enterprise customers:

Incident Level	First Response Time	Status Update Frequency
Critical	1 business day	Every 3 business days
Major	2 business days	Weekly
Minor	3 business days	As progress occurs

A “response” means acknowledgement + initial assessment.
Fix timelines are best-effort unless separately negotiated.

4. Support Channels

Enterprise support will be available through:

- Dedicated enterprise email (being finalized now)
- Optional pre-scheduled call windows (launching early 2026)
- Private issue tracking for Enterprise customers

Temporary contact instructions will be provided at purchase until automated systems are fully deployed.

5. SLA Coverage

Covered under this SLA:

- Installation and setup issues
 - Configuration guidance
 - Reproducible bugs
 - Integration questions
 - Clarification of expected behaviors
 - Priority handling of Critical issues
 - Access to hotfixes or patch releases (when applicable)
-

6. Exclusions

This SLA does not cover:

- Custom development
- Trading strategy design or modeling consulting
- Performance optimization of customer environments
- Issues caused by unsupported OS/hardware
- Failures originating from third-party libraries
- Customer-modified source code
- Operational uptime (this is packaged software, not a hosted SaaS)

These may be available under a separate consulting agreement.

7. Scheduled Maintenance

Fox ML Infrastructure may perform scheduled maintenance for:

- Security patches
- Dependency updates
- Infrastructure improvements
- Documentation revisions

Customers will be notified of substantial changes when feasible.

8. Availability Disclaimer

Fox ML Infrastructure is self-hosted software running in the customer's environment. As such, **no uptime or hosting guarantees** are provided under this SLA.

Customers are responsible for:

- Execution environment stability
 - Networking
 - Hardware reliability
 - Hosting and deployment
-

9. Limitation of Liability

Total liability under this SLA is limited to the fees paid for the Enterprise tier in the preceding 12 months. Fox ML Infrastructure is not liable for trading losses, financial impacts, downtime, or consequential damages.

10. SLA Evolution

This SLA may be updated as Fox ML Infrastructure matures.
Enterprise customers will be notified of material changes.
